



Lab 1 - DDPM



Task 1 - Swiss Roll

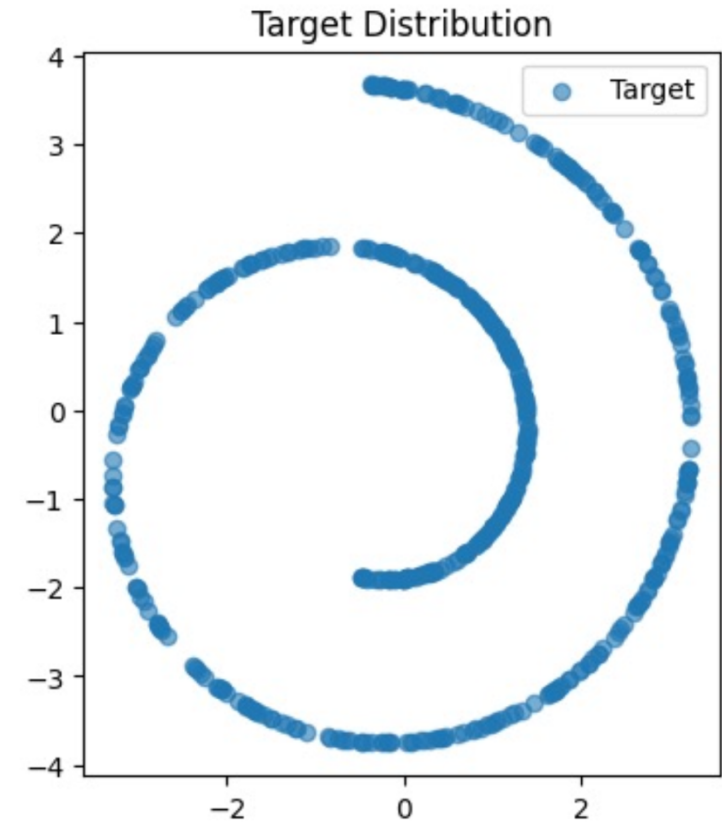
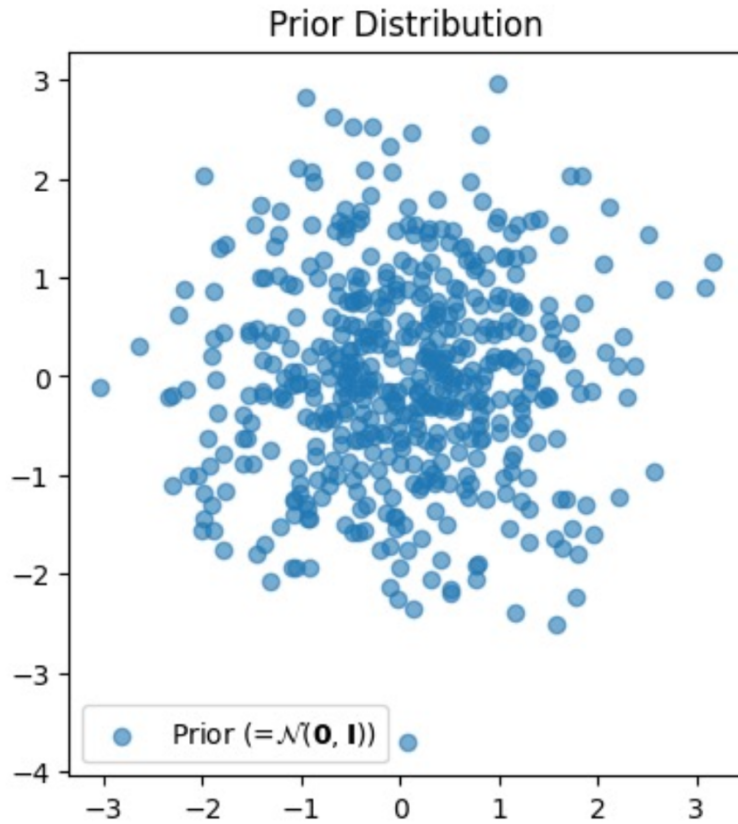


Task 2 - Image Generation

[DDPM paper](#)

Task 1 - Swiss Roll

Let's begin by modeling a simple distribution of 2D points ("Swiss Roll").



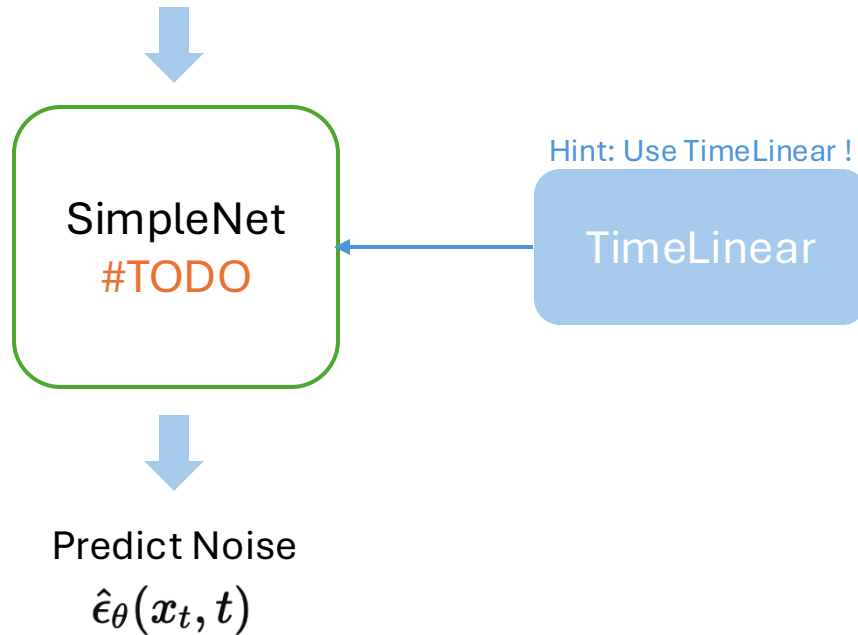
Task 1

#TODO1 - SimpleNet

- 2d_plot_diffusion_todo/network.py

Reverse Process

1. Noise data x_t
2. Current timestep t



```
class SimpleNet(nn.Module):
    def __init__(
        self, dim_in: int, dim_out: int, dim_hids: List[int], num_timesteps: int
    ):
        super().__init__()
        """
        (TODO) Build a noise estimating network.

        Args:
            dim_in: dimension of input
            dim_out: dimension of output
            dim_hids: dimensions of hidden features
            num_timesteps: number of timesteps
        """

        ##### TODO #####
        # DO NOT change the code outside this part.

        #####

    def forward(self, x: torch.Tensor, t: torch.Tensor):
        """
        (TODO) Implement the forward pass. This should output
        the noise prediction of the noisy input x at timestep t.

        Args:
            x: the noisy data after t period diffusion
            t: the time that the forward diffusion has been running
        """

        ##### TODO #####
        # DO NOT change the code outside this part.

        #####


        return x
```

Task 1

Forward Process

#TODO2 - q_sample

- 2d_plot_diffusion_todo/ddpm.py



$$q(x_t | x_0) = N(x_t; \sqrt{\bar{\alpha}_t} x_0, (1 - \bar{\alpha}_t) I).$$

```
def q_sample(self, x0, t, noise=None):
    """
    sample x_t from q(x_t | x_0) of DDPM.

    Input:
        x0 (`torch.Tensor`): clean data to be mapped to timestep t in the forward process of DDPM.
        t (`torch.Tensor`): timestep
        noise (`torch.Tensor`, optional): random Gaussian noise. if None, randomly sample Gaussian noise i
    Output:
        xt (`torch.Tensor`): noisy samples
    """
    if noise is None:
        noise = torch.randn_like(x0)

    ##### TODO #####
    # DO NOT change the code outside this part.
    # Compute xt.
    alphas_prod_t = extract(self.var_scheduler.alphas_cumprod, t, x0)
    xt = x0

    #####

    return xt
```

Task 1

Forward Process

#TODO2 - q_sample

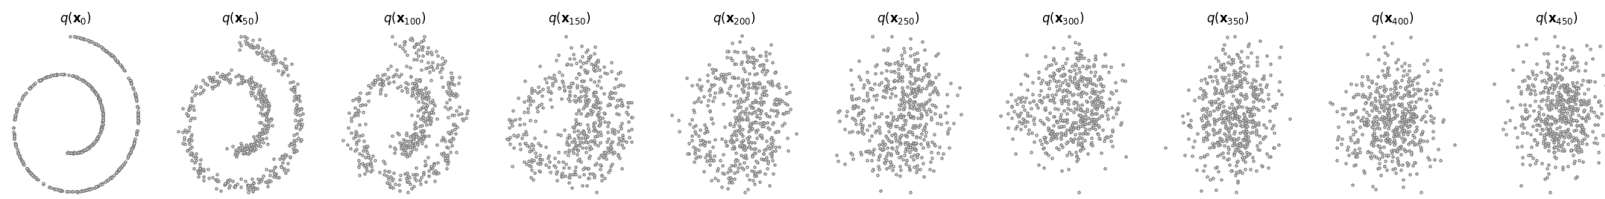
- 2d_plot_diffusion_todo/ddpm_tutorial.py

Check your implementation of q_sample !

Visualize $q(\mathbf{x}_t)$

```
fig, axs = plt.subplots(1, 10, figsize=(28, 3))
for i, t in enumerate(range(0, 500, 50)):
    x_t = ddpm.q_sample(target_ds[:num_vis_particles].to(device), (torch.ones(num_vis_particles, 1).to(device),))
    x_t = x_t.cpu()
    axs[i].scatter(x_t[:,0], x_t[:,1], color='white', edgecolor='gray', s=5)
    axs[i].set_axis_off()
    axs[i].set_title('$q(\mathbf{x}_{'+str(t)+'})$')
```

Python



Task 1

Reverse Process

#TODO3 - p_sample

- 2d_plot_diffusion_todo/ddpm.py

How to get $x_t \xrightarrow{p_\theta} x_{t-1}$?

1. Predict noise from SimpleNet

$$2. \mu_\theta(x_t, t) = \frac{1}{\sqrt{\alpha_t}} \left(x_t - \frac{\beta_t}{\sqrt{1-\bar{\alpha}_t}} \hat{\epsilon}_\theta(x_t, t) \right) \quad \text{Eq. 11}$$

$$3. \tilde{\beta}_t = \frac{1-\bar{\alpha}_{t-1}}{1-\bar{\alpha}_t} \beta_t \quad \text{Sec. 3.2}$$

$$4. x_{t-1} = \mu_\theta(x_t, t) + \sqrt{\tilde{\beta}_t} z, \quad z \sim \mathcal{N}(0, I) \quad \text{Eq. 11}$$

```
def p_sample(self, xt, t):
    """
    One step denoising function of DDPM:  $x_t \rightarrow x_{t-1}$ .
    Input:
    | xt (`torch.Tensor`): samples at arbitrary timestep t.
    | t (`torch.Tensor`): current timestep in a reverse process.
    Output:
    | x_t_prev (`torch.Tensor`): one step denoised sample. (=  $x_{t-1}$ )
    """
    ##### TODO #####
    # DO NOT change the code outside this part.
    # compute x_t_prev.
    if isinstance(t, int):
        t = torch.tensor([t]).to(self.device)
    eps_factor = (1 - extract(self.var_scheduler.alphas, t, xt)) / (
        1 - extract(self.var_scheduler.alphas_cumprod, t, xt)
    ).sqrt()

    beta_t = extract(self.var_scheduler.betas, t, xt) #  $\beta_t$ 
    alpha_t = extract(self.var_scheduler.alphas, t, xt) #  $\alpha_t$ 
    alpha_bar_t = extract(self.var_scheduler.alphas_cumprod, t, xt) #  $\bar{\alpha}_t$ 
    t_prev = (t - 1).clamp(min=0)
    alpha_bar_t_prev = extract(self.var_scheduler.alphas_cumprod, t_prev, xt) #  $\bar{\alpha}_{t-1}$ 

    # 1. predict noise
    # 2. Posterior mean
    # 3. Posterior variance
    # 4. Reverse step

    #####
    return x_t_prev
```

Task 1

Reverse Process

#TODO4 - p_sample loop

- 2d_plot_diffusion_todo/ddpm.py

```
def p_sample_loop(self, shape):
    """
    The loop of the reverse process of DDPM.

    Input:
    |   shape (`Tuple`): The shape of output. e.g., (num particles, 2)
    Output:
    |   x0_pred (`torch.Tensor`): The final denoised output through the DDPM reverse process.
    """
    ##### TODO #####
    # DO NOT change the code outside this part.
    # sample x0 based on Algorithm 2 of DDPM paper.
    x0_pred = torch.zeros(shape).to(self.device)

    #####
    return x0_pred
```


Task 1

#TODO5 - compute_loss

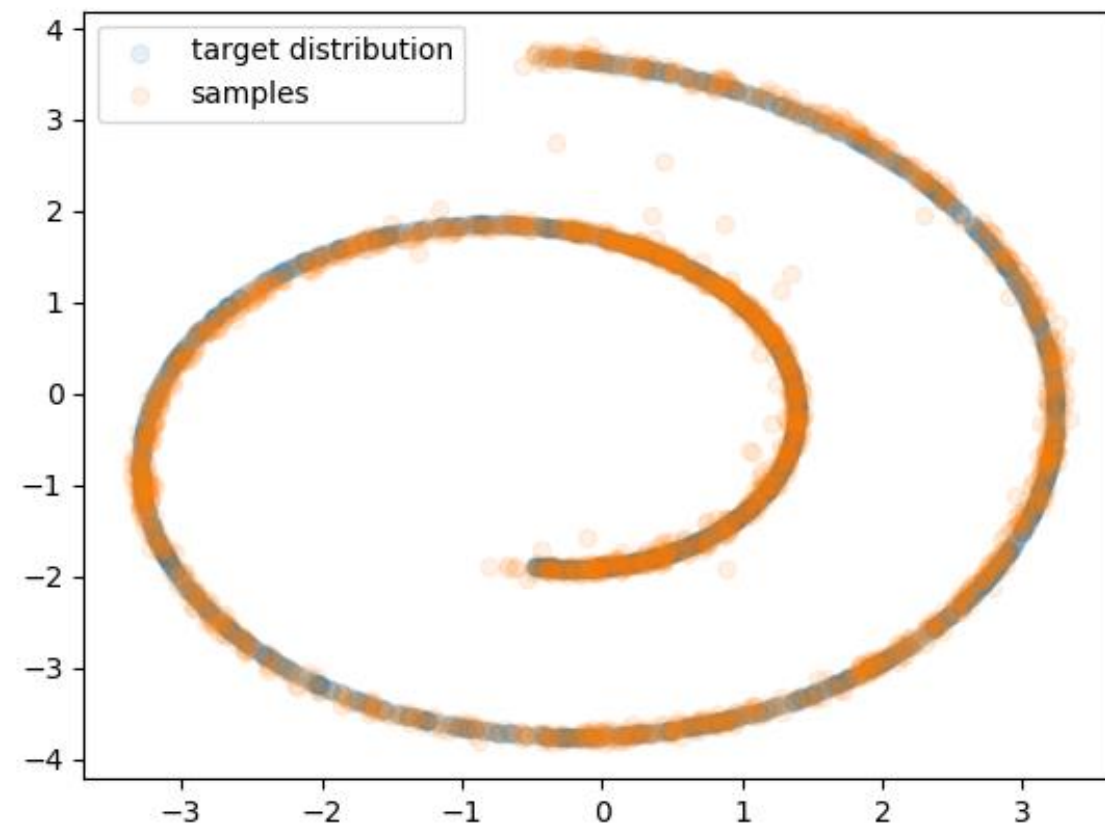
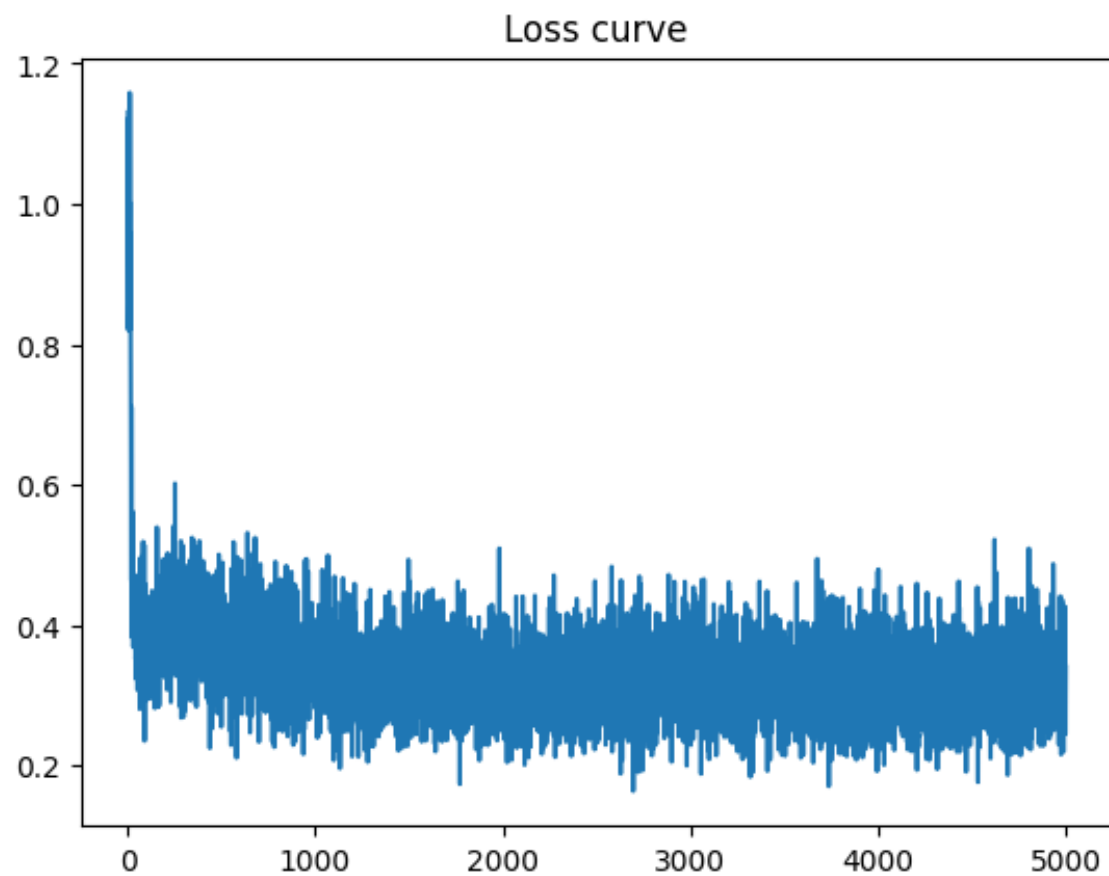
- 2d_plot_diffusion_todo/ddpm.py

```
def compute_loss(self, x0):  
    """  
    The simplified noise matching loss corresponding Equation 14 in DDPM paper.  
    Input:  
    |   x0 (`torch.Tensor`): clean data  
    Output:  
    |   loss: the computed loss to be backpropagated.  
    """  
  
    ##### TODO #####  
    # DO NOT change the code outside this part.  
    # compute noise matching loss.  
    batch_size = x0.shape[0]  
  
    # 1) random choose timestep  
    t = (  
        torch.randint(0, self.var_scheduler.num_train_timesteps, size=(batch_size,))  
        .to(x0.device)  
        .long()  
    )  
    # 2) get GT noise, and use q_sample to get x_t  
  
    # 3) predict noise  
  
    # 4) MSE loss (eps, eps_pred)  
  
    loss = None  
  
    #####  
    return loss
```


Task 1

#TODO5 – train & evaluate


- 2d_plot_diffusion_todo/ddpm_tutorial.py



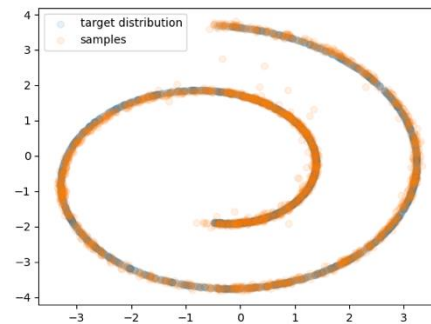
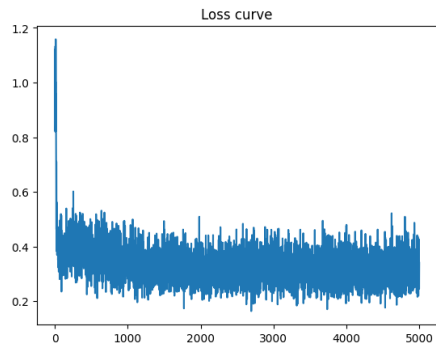
Task 1

#Report

1. Complete and Explain of all **#TODO** code (20pts)
2. Fig of your implementation of `q_sample` (10pts)



3. Fig of loss curve and evaluation result (10 pts)



Task 2 – Image Generation



What to Do: Task 2

*Ugh...
Time to write more code...*



Samples from our model trained using the [AFHQ dataset](#).

Bring your codes from Task 1!

The code needs to be modified, but the changes should be kept minimal.

- `q_sample` → `add_noise`;
- `p_sample` → `step`;
- `compute_loss` → `get_loss`.

Task 2

Forward Process

TODO1 – add_noise

- image_diffusion_todo/scheduler.py

```
def add_noise(
    self,
    x_0: torch.Tensor,
    t: torch.IntTensor,
    eps: Optional[torch.Tensor] = None,
):
    """
    A forward pass of a Markov chain, i.e.,  $q(x_t | x_0)$ .

    Input:
        x_0 (`torch.Tensor [B,C,H,W]`): samples from a real data distribution  $q(x_0)$ .
        t: (`torch.IntTensor [B]`)
        eps: (`torch.Tensor [B,C,H,W]`, optional): if None, randomly sample Gaussian noise in the function.
    Output:
        x_t: (`torch.Tensor [B,C,H,W]`): noisy samples at timestep t.
        eps: (`torch.Tensor [B,C,H,W]`): injected noise.
    """

    if eps is None:
        eps = torch.randn(x_0.shape, device='cuda')

    ##### TODO #####
    # DO NOT change the code outside this part.
    # Assignment 1. Implement the DDPM forward step.
    x_t = None

    #####

    return x_t, eps
```

Hint : Refer to 2d_plot_diffusion_todo/ddpm.py – **q_sample**

Task 2

TODO4 – beta scheduling

- image_diffusion_todo/scheduler.py

We know

$$q(x_t | x_0) = \sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, \quad \epsilon \sim \mathcal{N}(0, I)$$

and

$$\alpha_t = 1 - \beta_t$$

$$\bar{\alpha}_t = \prod_{i=1}^t \alpha_i$$

Implement the **cosine scheduler**.

```
if mode == "linear":
    betas = torch.linspace(beta_1, beta_T, steps=num_train_timesteps)
elif mode == "quad":
    betas = (
        torch.linspace(beta_1**0.5, beta_T**0.5, num_train_timesteps) ** 2
    )
elif mode == "cosine":
    ##### TODO #####
    # Implement the cosine beta schedule (Nichol & Dhariwal, 2021).
    # Hint:
    # 1. Define alphā_t = f(t/T) where f is a cosine schedule:
    #     alphā_t = cos^2( ( t/T + s ) / (1+s) ) * (π/2) )
    #     with s = 0.008 (a small constant for stability).
    # 2. Convert alphā_t into betas using:
    #     beta_t = 1 - alphā_t / alphā_{t-1}
    # 3. Return betas as a tensor of shape [num_train_timesteps].
    raise NotImplementedError("TODO: Implement cosine beta schedule here!")
```


Task 2

TODO5 – predictor

- image_diffusion_todo/scheduler.py

Implement the **reverse step** according to the chosen **predictor** (**noise**, **x_0** , or **mean**).

```
def step_predict_noise(self, x_t: torch.Tensor, t: int, eps_theta: torch.Tensor):
    """
    Noise prediction version (the standard DDPM formulation).

    Input:
    | x_t: noisy image at timestep t
    | t: current timestep
    | eps_theta: predicted noise  $\hat{\epsilon}_\theta(x_t, t)$ 
    Output:
    | sample_prev: denoised image sample at timestep t-1
    """
    ##### TODO #####
    # 1. Extract beta_t, alpha_t, and alpha_bar_t from the scheduler.
    # 2. Compute the predicted mean  $\mu_\theta(x_t, t) = 1/\sqrt{\alpha_t} * (x_t - (\beta_t/\sqrt{1-\alpha_t})) * \epsilon_\theta(x_t, t)$ 
    # 3. Compute the posterior variance  $\tilde{\beta}_t = ((1-\alpha_{t-1}))/((1-\alpha_t)) * \beta_t$ .
    # 4. Add Gaussian noise scaled by  $\sqrt{(\tilde{\beta}_t)}$  unless t == 0.
    # 5. Return the final sample at t-1.
    sample_prev = None
    #####
    return sample_prev
```

```
def step_predict_x0(self, x_t: torch.Tensor, t: int, x0_pred: torch.Tensor):
    """
    x0 prediction version (alternative DDPM objective).

    Input:
    | x_t: noisy image at timestep t
    | t: current timestep
    | x0_pred: predicted clean image  $\hat{x}_0(x_t, t)$ 
    Output:
    | sample_prev: denoised image sample at timestep t-1
    """
    ##### TODO #####

    sample_prev = None
    #####
    return sample_prev

def step_predict_mean(self, x_t: torch.Tensor, t: int, mean_theta: torch.Tensor):
    """
    Mean prediction version (directly outputting the posterior mean).

    Input:
    | x_t: noisy image at timestep t
    | t: current timestep
    | mean_theta: network-predicted posterior mean  $\hat{\mu}_\theta(x_t, t)$ 
    Output:
    | sample_prev: denoised image sample at timestep t-1
    """
    ##### TODO #####

    sample_prev = None
    #####
    return sample_prev
```


Task 2

TODO5 – predictor

- image_diffusion_todo/model.py

Implement **loss functions** for training the network to predict different targets: **noise**, **x_0** , or the **posterior mean**.

```
def get_loss(self, x0, class_label=None, noise=None):
    if self.predictor == "noise":
        return self.get_loss_noise(x0, class_label, noise)
    elif self.predictor == "x0": ##### TODO
        return self.get_loss_x0(x0, class_label, noise)
    elif self.predictor == "mean": ##### TODO
        return self.get_loss_mean(x0, class_label, noise)
    else:
        raise ValueError(f"Unknown predictor: {self.predictor}")
```

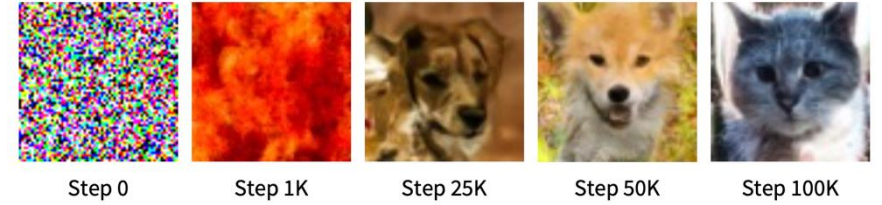
```
def get_loss_x0(self, x0, class_label=None, noise=None):
    ##### TODO #####
    # Here we implement the "predict x0" version.
    # 1. Sample a timestep and add noise to get (x_t, noise).
    # 2. Pass (x_t, timestep) into self.network, where the output should represent the clean sample x0_pred.
    # 3. Compute the loss as MSE(predicted x0_pred, ground-truth x0).
    #####
    loss = None
    return loss

def get_loss_mean(self, x0, class_label=None, noise=None):
    ##### TODO #####
    # Here we implement the "predict mean" version.
    # 1. Sample a timestep and add noise to get (x_t, noise).
    # 2. Pass (x_t, timestep) into self.network, where the output should represent the posterior mean  $\mu_\theta(x_t, t)$ .
    # 3. Compute the *true* posterior mean from the closed-form DDPM formula (using x0, x_t, noise, and schedule).
    # 4. Compute the loss as MSE(predicted mean, true mean).
    #####
    loss = None
    return loss
```

Task 2

TODO6 - train

- image_diffusion_todo/train.py



```
python train.py --mode {BETA_SCHEDULING} --predictor {PREDICTOR}
```

Try different **beta scheduling**:

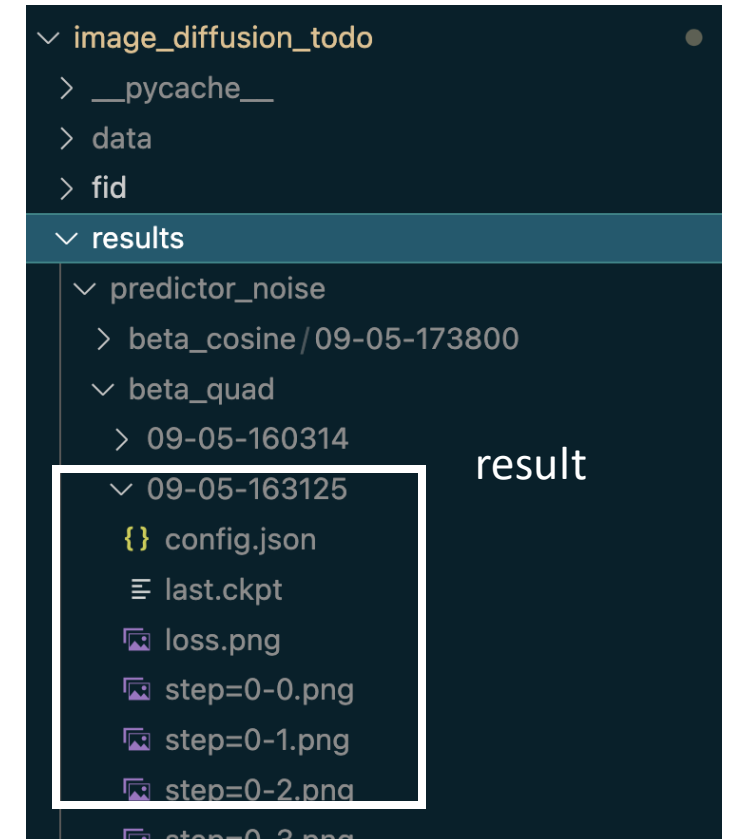
1. linear
2. quadratic
3. cosine

Try different **predictor**:

1. noise (default)
2. mean
3. x_0

Please train for at least **50,000–100,000** iterations.

Training may take **over 6 hours per time**, so start ASAP!



Task 2

TODO7 - sampling

- image_diffusion_todo/sampling.py

```
python sampling.py --ckpt_path {CKPT} --save_dir {SAMPLING_SAVE_DIR} \  
--mode {BETA_SCHEDULING} --predictor {PREDICTOR}
```

▼ predictor_x0 / beta_linear

> 09-05-201510

> 09-05-202134

> 09-05-204130

▼ 09-06-022830

{ } config.json

≡ last.ckpt

🖼️ loss.png

🖼️ step=0-0.png

🖼️ step=0-1.png

Task 2

TODO8 - evaluate

- image_diffusion_todo/dataset.py

Prepare the data for evaluation by running

python dataset.py (Only once!)

This will create the eval directory under data/afhq.

Do NOT forget to run this. Otherwise, you will get incorrect FIDs!

cd fid

python measure_fid.py ../data/afhq/eval {SAMPLING_SAVE_DIR}

FID: 229.33834138594412

Incorrect FIDs

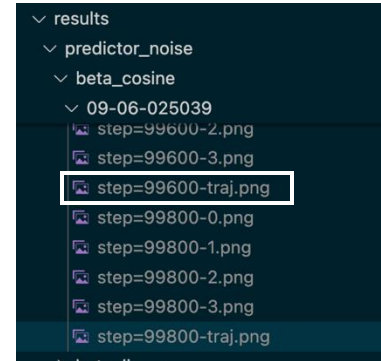
FID: 10.844529455220403

Correct FIDs

Task 2

#Report

1. Complete and Explain of all **#TODO** code. (30pts)
2. Show the **trajectory fig** and compare results across the three **beta schedulers**.
(linear, quadratic, cosine) (10pts)



3. Show the results of different **predictors** (noise, x_0 , posterior mean) and discuss. (10pts)



8 samples per predictor, 24 in total.

4. Screenshot of the **Best FID** of your training result, explain the training setting. (20pts)
 - 20 pts: $FID < 15$
 - 15 pts: $15 \leq FID < 20$
 - 10 pts: $20 \leq FID < 30$
 - 5 pts: $30 \leq FID < 40$
 - 0 pts: $FID \geq 40$

Submit

Create a single ZIP file {ID}_lab1.zip including:

- The **PDF** file formatted following the guideline;
- Your **code** without checkpoints for DDPMs and the Inception Network

412551014_lab1.zip

- report.pdf
- 2d_plot_diffusion_todo
- image_diffusion_todo



Thank you